

Comparison between Intel® DPDK's OpenSSL crypto and ZLIB compression Poll Mode Drivers

Written by Gereltsetseg Altangerel/Ph.D. student/
Supervised by Tejfel Máté/Ph.D./

ELTE Eötvös Loránd University, Budapest, Hungary
Faculty of Informatics, 3in Research Group, Martonvásár, Hungary

¹ The research has been supported by the European Union, co-financed by the European Social Fund (EFOP-3.6.2-16-2017-00013, Thematic Fundamental Research Collaborations Grounding Innovation in Informatics and Infocommunications).

Table of Contents

1. Introduction	3
2. Intel® DPDK overview	3
2.1 Environment Abstraction Layer	3
2.2 Managing Queues: rte_ring	4
2.3 Memory Management: rte_mempool.....	4
2.4 Buffer Management: rte_mbuf.....	4
3. DPDK Poll mode drivers (PMDs)	5
4. Compression Zlib and crypto OpenSSL PMDs.....	5
4.1 Compression operation representation	6
4.2 Crypto operation representation	6
4.3 How compression and crypto operation to be processed?	6
4.4 Comparison of Crypto OpenSSL and Zlib Compression PMDs	7
5. Compression application	20
6. Conclusion	21
ACKNOWLEDGMENT	21
REFERENCE	21
APPENDIX 1: DPDK, crypto OpenSSL PMD, compression ZLIB PMD installation guide.....	22
APPENDIX 2: Source code of application using ZLIB compression PMD.....	22

1. Introduction

Nowadays packet processing is implemented in software systems such as the Unix operating system instead of dedicated hardware. Although this solution gives large flexibility (deploying features rapidly in software) and cost-effectiveness, it limits packet processing speed. Main factors for limiting throughput are CPU bottleneck and usage of main memory. The latter one is divided into three problems: 1. per packet memory allocation and deallocation 2. complex `sk_buff` data structure 3. multiple memory copies. [1]. In order to solve these problems, high-speed packet processing frameworks such as Data Plane Development Kit (DPDK), netmap, and PF_RING are developed by using commodity hardware. These frameworks replace or extend conventional concepts by implementing their own driver and kernel level improvements. Moreover, the Intel® DPDK framework provides a set of software libraries and drivers for fast packet processing in dataplane applications of the Intel architectures [2, 3, 5]. In this technical report, based on application using compress ZLIB and crypto OpenSSL Poll mode drivers(PMDs) of DPDK, we compared these two PMDs in terms of libraries and functions.

2. Intel® DPDK overview

Figure 1 shows the DPDK components: EAL(Environment Abstraction Layer) and several libraries that are optimized for high performance as mentioned above. First of all, DPDK fulfills basic tasks, similar to what the Linux network stack does: allocating memory for network packets, buffering packet descriptors in ring-like structures and passing the packets from the NIC to the application (and vice versa). We are presenting EAL, and libraries that are necessary for these assignments include memory, queue and buffer management in the following sections briefly..

2.1 Environment Abstraction Layer (EAL)

The EAL is the main concept behind the DPDK. The EAL is a set of programming tools that let the DPDK work in a specific hardware environment and under a specific operating system. In the official DPDK repository, libraries and drivers that are part of the EAL are saved in the `rte_eal` directory. [6] Moreover, EAL is responsible for gaining access to low-level resources such as hardware and memory space. Also, it provides a generic interface that hides the environment specifics from the applications and libraries and is the responsibility of the initialization routine to decide how to allocate these resources (memory space, devices, timers, consoles, and so on).[5] In other words, the EAL performs all works of the kernel such as loading, launching applications, and reserving memory.[2,5] Please look at figure 2 to see the difference between packet processing with DPDK and without DPDK.

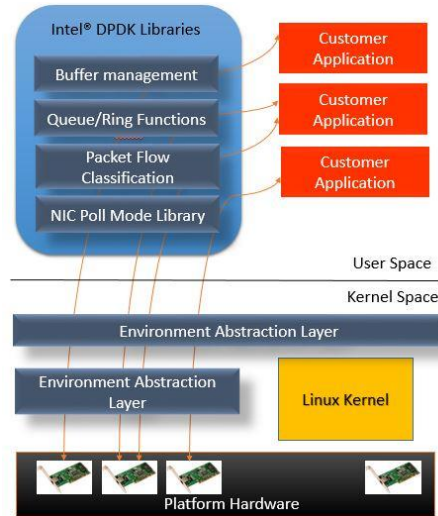


Figure 1: Intel® DPDK EAL and libraries

2.2 Managing Queues: `rte_ring`

In the traditional packet processing, packets received by the network card are sent to a ring buffer, which acts as a receiving queue. Packets received in the DPDK are also sent to a queue implemented on the `rte_ring` library. More information about the library's description could be found from the developer's guide and comments in the source code.[5]

2.3 Memory Management: `rte_mempool`

The EAL provides a mapping of physical memory [2]. It creates a table of descriptors which are called `rte_memseg` which each point to a contiguous portion of memory to circumvent the problem that available physical memory can have gaps. Those segments can then be divided into memory zones. These zones are the basic memory unit used by any further object created from an application or other libraries. One of these objects is the `rte_mempool` provided by the `lib_rte_mempool` library.

2.4 Buffer Management: `rte_mbuf`

In the Linux network stack, all network packets are represented by the `sk_buff` data structure. In DPDK, this is done using the `rte_mbuf` struct, which is provided by the `rte_mbuf` library. The buffer management approach in DPDK is reminiscent of the approach used in FreeBSD: instead of one big `sk_buff` struct, there are many smaller `rte_mbuf` buffers. The buffers are created before the DPDK application is launched and are stored in a mempool (memory is allocated by `rte_mempool`).

3. DPDK Poll mode drivers (PMDs)

In Figure 2, the left side shows the traditional way of packet processing on Linux, and the right side shows the packet processing way with DPDK. The traditional packet processing is very complicated and leads to some problems such as consuming a lot of bus cycles, context switching, interrupts and so on.

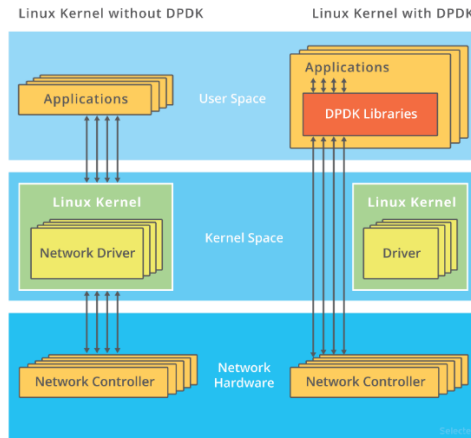


Figure 2: Linux kernel with and without DPDK

As we can see, the kernel in the right side is being skipped: interactions with the network controller are performed via special drivers and libraries. These special drivers are the DPDK poll mode drivers (PMD) which organize all further communication between the user application using DPDK libraries and network card. More specifically, a PMD accesses the RX and TX descriptors directly without any interrupts (with the exception of Link Status Change interrupts) to quickly receive, process and deliver packets in the user's application.[2] Furthermore, DPDK has developed poll mode drivers for all supported network cards and virtual devices.

4. Compression Zlib and crypto OpenSSL PMDs

A Poll Mode Driver (PMD) consists of APIs to configure the devices and their respective queues.

OpenSSL Crypto PMD: The cryptodev library provides a Crypto device framework for management and provisioning of hardware and software Crypto poll mode drivers, defining generic APIs which support a number of different Crypto operations. The framework currently only supports cipher, authentication, chained cipher/authentication and AEAD symmetric and asymmetric Crypto operations.[2]

Zlib Compress PMD: The compression framework provides a generic set of APIs to perform compression services as well as to query and configure compression devices both physical(hardware) and virtual(software) to perform those services. The framework currently only supports lossless compression schemes: Deflate and LZS.[2]

4.1 Compression operation representation

Compression operation is described via struct `rte_comp_op`, which contains both input and output data. The operation structure includes the operation type (stateless or stateful), the operation status and the `priv_xform/stream` handle, source, destination and checksum buffer pointers. It also contains the source mempool from which the operation is allocated. PMD updates consumed field with amount of data read from source buffer and produced field with amount of data of written into destination buffer along with status of operation. See section Produced, Consumed And Operation Status for more details.

Compression operations mempool also has an ability to allocate private memory with the operation for application's purposes. Application software is responsible for specifying all the operation specific fields in the `rte_comp_op` structure which are then used by the compression PMD to process the requested operation. [2]

4.2 Crypto operation representation

An Crypto operation is represented by an `rte_crypto_op` structure, which is a generic metadata container for all necessary information required for the Crypto operation to be processed on a particular Crypto device poll mode driver. The operation structure includes the operation type, the operation status and the session type (session-based/less), a reference to the operation specific data, which can vary in size and content depending on the operation being provisioned. It also contains the source mempool for the operation, if it allocated from a mempool.

If Crypto operations are allocated from a Crypto operation mempool, see next section, there is also the ability to allocate private memory with the operation for applications purposes.

Application software is responsible for specifying all the operation specific fields in the `rte_crypto_op` structure which are then used by the Crypto PMD to process the requested operation.[2]

4.3 How compression and crypto operation to be processed?

There are two kinds of crypto or compression devices: physical and virtual device. A queue pair on a compress device accepts a burst of compress operations using enqueue burst API. On physical compress devices, the enqueue burst API will place the operations to be processed on the input queue of a device's hardware. In terms of virtual devices, the compress operations' processing is usually completed when the enqueue is invoked to the compressing device. The dequeue burst API will retrieve any processed operations available from the queue pair on the compress device, from physical devices this is usually directly from the devices processed queue, and for virtual device's from a `rte_ring` where processed operations are place after being processed on the enqueue call. This procedure is same in crypto operation. If you want to get more information about crypto and compression operations descriptions, please access to official references.[2]

4.4 Comparison of Crypto OpenSSL and Zlib Compression PMDs

Appendix 2 is the source code for compression application. This code was tested in Ubuntu 18.04.2, DPDK 19.02 and 19.05. Refer to Appendix 1 for DPDK and compression PMD installation instructions.

All functions and structures associated with the source code are explained in the following reference table. For example, at the end of the table, the "Line 29" mark means that function or structure in Line 29 of the source code is explained in this cell of the table.

	Crypto PMD	Compression PMD
1	int rte_eal_init (int argc, char ** argv) - Initialize the Environment Abstraction Layer (EAL). This function is to be executed on the MASTER lcore only, as soon as possible in the application's main() function. The function finishes the initialization process before main() is called. It puts the SLAVE lcores in the WAIT state. Parameters argc A non-negative value. If it is greater than 0, the array members for argv[0] through argv[argc] (non-inclusive) shall contain pointers to strings. argv An array of strings. The contents of the array, as well as the strings which are pointed to by the array, may be modified by this function. Returns -On success, the number of parsed arguments, which is greater or equal to zero. After the call to rte_eal_init(), all arguments argv[x] with x < ret may have been modified by this function call and should not be further interpreted by the application. The EAL does not take any ownership of the memory used for either the argv array, or its members. -On failure, -1 and rte_errno is set to a value indicating the cause for failure. In some instances, the application will need to be restarted as part of clearing the issue. <div style="text-align: right;">"line 29"</div>	
	The RTE mempool structure. (#include <rte_mempool.h>, this header is included in rte_compressed.h) rte_mempool – data structure for creating packet memory buffers and operation pools. Struct Data Fields char name [RTE_MEMZONE_NAMESIZE] – Name of mempool void * pool_config - Ring or pool to store objects. struct rte_memzone * mz unsigned int flags int socket_id uint32_t size uint32_t cache_size uint32_t elt_size uint32_t header_size uint32_t trailer_size unsigned private_data_size int32_t ops_index struct rte_mempool_cache * local_cache uint32_t populated_size struct rte_mempool_objhdr_list elt_list uint32_t nb_mem_chunks struct rte_mempool_memhdr_list mem_list void * pool_data	

	uint64_t pool_id			"line 33"
2	rte_pktmbuf_pool_create (const char * name, unsigned n, unsigned cache_size, uint16_t priv_size, uint16_t data_room_size, int socket_id) - Create a mbuf pool. This function creates and initializes a packet mbuf pool. It is a wrapper to rte_mempool functions. Parameters name The name of the mbuf pool. n The number of elements in the mbuf pool. The optimum size (in terms of memory usage) for a mempool is when n is a power of two minus one: n = (2^q - 1). cache_size Size of the per-core object cache. See rte_mempool_create() for details. priv_size Size of application private area between the rte_mbuf structure and the data buffer. This value must be aligned to RTE_MBUF_PRIV_ALIGN. data_room_size Size of data buffer in each mbuf, including RTE_PKTMBUF_HEADROOM. socket_id The socket identifier where the memory should be allocated. The value can be SOCKET_ID_ANY if there is no NUMA constraint for the reserved zone. socket_id The socket identifier where the memory should be allocated. The value can be SOCKET_ID_ANY if there is no NUMA constraint for the reserved zone Returns The pointer to the new allocated mempool, on success. NULL on error with rte_errno set appropriately. Possible rte_errno values include: -E_RTE_NO_CONFIG - function could not get pointer to rte_config structure -E_RTE_SECONDARY - function was called from a secondary process instance -EINVAL - cache size provided is too large, or priv_size is not aligned. -ENOSPC - the maximum number of memzones has already been allocated -EEXIST - a memzone with the same name already exists -ENOMEM - no appropriate memory area found in which to create memzone			
	"line 36"			
	struct rte_mempool* rte_crypto_op_pool_create (const char * name, enum rte_crypto_op_type type, unsigned nb_elts, unsigned cache_size, uint16_t priv_size, int socket_id) -Creates a crypto operation pool Parameters name pool name type crypto operation type, use RTE_CRYPTO_OP_TYPE_UNDEFINED for a pool which supports all operation types nb_elts number of elements in pool cache_size Number of elements to cache on lcore, see rte_mempool_create for further details about cache size priv_size Size of private data to allocate with each operation socket_id Socket to allocate memory on Returns On success pointer to mempool On failure NULL	struct rte_mempool* __rte_experimental rte_comp_op_pool_create (const char * name, unsigned int nb_elts, unsigned int cache_size, uint16_t user_size, int socket_id)-Creates an operation pool Parameters name Compress pool name nb_elts Number of elements in pool cache_size Number of elements to cache on lcore, see rte_mempool_create for further details about cache size user_size Size of private data to allocate for user with each operation socket_id Socket to identifier allocate memory on Returns On success pointer to mempool On failure NULL		"line 47"
	struct rte_mempool* __rte_experimental rte_cryptodev_sym_session_pool_create (const char * name, uint32_t nb_elts, uint32_t elt_size, uint32_t cache_size, uint16_t priv_size, int socket_id)-Create a symmetric session mempool.			

	<p>Parameters name The unique mempool name. nb_elts The number of elements in the mempool. elt_size The size of the element. This value will be ignored if it is smaller than the minimum session header size required for the system. For the user who want to use the same mempool for sym session and session private data it can be the maximum value of all existing devices' private data and session header sizes. cache_size The number of per-lcore cache elements priv_size The private data size of each session. socket_id The socket_id argument is the socket identifier in the case of NUMA. The value can be SOCKET_ID_ANY if there is no NUMA constraint for the reserved zone. Returns On success return size of the session On failure returns 0</p>	
3	<p>int rte_vdev_init (const char * name, const char * args) - Initialize a driver specified by name. (Create compress, decompress, encrypt and decrypt device)</p> <p>Parameters name The pointer to a driver name to be initialized. args The pointer to arguments used by driver initialization.</p> <p>Returns 0 on success, negative on error</p> <p style="text-align: right;">”line 59, 67”</p>	
	<p>struct rte_cryptodev_config - Crypto device configuration structure is used to pass the configuration parameters for socket selection and number of queue pairs.</p> <p>Data Fields int socket_id Socket to allocate resources on uint16_t nb_queue_pairs Number of queue pairs to configure on device</p>	<p>rte_compressdev_config - structure is used to pass configuration parameters.</p> <p>Data Fields int socket_id Socket on which to allocate resources uint16_t nb_queue_pairs Total number of queue pairs to configure on a device uint16_t max_nb_priv_xforms Max number of private_xforms which will be created on the device uint16_t max_nb_streams Max number of streams which will be created on the device</p> <p style="text-align: right;">”line 73”</p>
	<p>rte_cryptodev_qp_conf - Crypto device queue pair configuration structure.</p> <p>Data Fields uint32_t nb_descriptors Number of descriptors per queue pair struct rte_mempool * mp_session The mempool for creating session in sessionless mode struct rte_mempool * mp_session_private The mempool for creating sess private data in sessionless mode</p>	

Device configuration	
<p>int rte_cryptodev_configure (uint8_t dev_id, struct rte_cryptodev_config *config) - Configure a device. This function must be invoked first before any other function in the API. This function can also be re-invoked when a device is in the stopped state.</p> <p>Parameters dev_id The identifier of the device to configure. config The crypto device configuration structure.</p> <p>Returns 0: Success, device configured. <0: Error code returned by the driver configuration function.</p>	<p>int __rte_experimental rte_compressdev_configure (uint8_t dev_id, struct rte_compressdev_config *config) Configure a device. This function must be invoked first before any other function in the API. This function can also be re-invoked when a device is in the stopped state.</p> <p>Parameters dev_id Compress device identifier config The compress device configuration</p> <p>Returns 0: Success, device configured. <0: Error code returned by the driver configuration function.</p> <p style="text-align: right;">”line 80,83,100,103”</p>
Configuration of Queue Pairs	
<p>int rte_cryptodev_queue_pair_setup (uint8_t dev_id, uint16_t queue_pair_id, const struct rte_cryptodev_qp_conf *qp_conf, int socket_id) - Allocate and set up a receive queue pair for a device.</p> <p>Parameters dev_id The identifier of the device. queue_pair_id The index of the queue pairs to set up. The value must be in the range [0, nb_queue_pair - 1] previously supplied to rte_cryptodev_configure(). qp_conf The pointer to the configuration data to be used for the queue pair. socket_id The socket_id argument is the socket identifier in case of NUMA. The value can be SOCKET_ID_ANY if there is no NUMA constraint for the DMA memory allocated for the receive queue pair.</p> <p>Returns 0: Success, queue pair correctly set up. <0: Queue pair configuration failed</p>	<p>int __rte_experimental rte_compressdev_queue_pair_setup (uint8_t dev_id, uint16_t queue_pair_id, uint32_t max_inflight_ops, int socket_id) - Allocate and set up a receive queue pair for a compress and decompress device. This should only be called when the device is stopped.</p> <p>Parameters dev_id Compress device identifier queue_pair_id The index of the queue pairs to set up. The value must be in the range [0, nb_queue_pair - 1] previously supplied to rte_compressdev_configure() max_inflight_ops Max number of ops which the qp will have to accommodate simultaneously socket_id The socket_id argument is the socket identifier in case of NUMA. The value can be SOCKET_ID_ANY if there is no NUMA constraint for the DMA memory allocated for the receive queue pair</p> <p>Returns 0: Success, queue pair correctly set up. <0: Queue pair configuration failed</p> <p style="text-align: right;">”line 86, 90, 106, 110”</p>
Start device	
<p>int rte_cryptodev_start (uint8_t dev_id) - Start an device.</p> <p>The device start step is the last one and consists of setting the configured offload features and in starting the transmit and the receive units of the device. On success, all basic functions exported by the API (link status, receive/transmit, and so on) can be invoked.</p> <p>Parameters</p>	<p>int __rte_experimental rte_compressdev_start (uint8_t dev_id) - Start a device.</p> <p>The device start step is called after configuring the device and setting up its queue pairs. On success, data-path functions exported by the API (enqueue/dequeue, etc) can be invoked.</p> <p>Parameters dev_id Compress device identifier</p> <p>Returns</p>

<p>dev_id The identifier of the device.</p> <p>Returns 0: Success, device started. <0: Error code of the driver device start function.</p>	<p>0: Success, device started. <0: Error code of the driver device start function. "line 94, 114"</p>
<p>rte_crypto_sym_xform-Symmetric crypto transform structure.</p> <p>This is used to specify the crypto transforms required, multiple transforms can be chained together to specify a chain transforms such as authentication then cipher, or cipher then authentication. Each transform structure can hold a single transform, the type field is used to specify which transform is contained within the union</p> <p>Data Fields struct rte_crypto_sym_xform * next next xform in chain enum rte_crypto_sym_xform_type type xform type struct rte_crypto_auth_xform auth Authentication / hash xform struct rte_crypto_cipher_xform cipher Cipher xform struct rte_crypto_aead_xform aead AEAD xform</p>	<p>rte_comp_xform Struct Reference</p> <p>Compression transform structure: This is used to specify the compression transforms required. Each transform structure can hold a single transform, the type field is used to specify which transform is contained within the union.</p> <p>enum rte_comp_xform_type type struct rte_comp_compress_xform compress struct rte_comp_decompress_xform decompress</p> <p>Data Fields enum rte_comp_xform_type type Compression transform types Enumerator: RTE_COMP_COMPRESS Compression service - compress RTE_COMP_DECOMPRESS Compression service – decompress</p> <p>struct rte_comp_compress_xform compress xform for compress operation</p> <p>rte_comp_compress_xform struct reference</p> <p>enum rte_comp_algorithm algo union { struct rte_comp_deflate_params deflate }; int level uint8_t window_size enum rte_comp_checksum_type checksum enum rte_comp_hash_algorithm hash_algo</p> <p>Compression Algorithms (algo) Enumerator: -RTE_COMP_ALGO_NULL No Compression algorithm No compression. Pass-through, data is copied unchanged from source buffer to destination buffer. -RTE_COMP_ALGO_DEFLATE DEFLATE compression algorithm https://tools.ietf.org/html/rfc1951 -RTE_COMP_ALGO_LZS LZS compression algorithm https://tools.ietf.org/html/rfc2395</p> <p>Compression Huffman Type - used by DEFLATE algorithm</p>

		<p>Enumerator:</p> <ul style="list-style-type: none"> -RTE_COMP_HUFFMAN_DEFAULT PMD may choose which Huffman codes to use -RTE_COMP_HUFFMAN_FIXED Use Fixed Huffman codes -RTE_COMP_HUFFMAN_DYNAMIC Use Dynamic Huffman codes <p>uint8_t window_size Base two log value of sliding window to be used. If window size can't be supported by the PMD then it may fall back to a smaller size. This is likely to result in a worse compression ratio.</p> <p>Compression Hash Algorithms</p> <p>Enumerator:</p> <ul style="list-style-type: none"> -RTE_COMP_HASH_ALGO_NONE No hash -RTE_COMP_HASH_ALGO_SHA1 SHA1 hash algorithm -RTE_COMP_HASH_ALGO_SHA2_256 SHA256 hash algorithm of SHA2 family <p>struct rte_comp_decompress_xform decompress</p> <p>Data Fields enum rte_comp_algorithm algo enum rte_comp_checksum_type chksum uint8_t window_size enum rte_comp_hash_algorithm hash_algo</p> <p>Note: All of the description for decompression are same with compression</p> <p style="text-align: right;">”line 120, 135”</p>
	<p>rte_cryptodev_sym_session struct rte_cryptodev_sym_session * rte_cryptodev_sym_session_create (struct rte_mempool *mempool) - Create symmetric crypto session header (generic with no private data)</p> <p>Parameters mempool Symmetric session mempool to allocate session objects from</p> <p>Returns On success return pointer to sym-session On failure returns NULL</p>	
	<p>rte_cryptodev_info - Crypto device information</p> <p>const char * driver_name uint8_t driver_id struct rte_device * device uint64_t feature_flags struct rte_cryptodev_capabilities * capabilities unsigned max_nb_queue_pairs</p>	<p>rte_compressdev_info - comp device information</p> <p>const char * driver_name uint64_t feature_flags struct rte_compressdev_capabilities * capabilities uint16_t max_nb_queue_pairs</p>

uint16_t min_mbuf_headroom_req uint16_t min_mbuf_tailroom_req unsigned max_nb_sessions	"line 144, 158"
void rte_cryptodev_info_get (uint8_t dev_id, struct rte_cryptodev_info * dev_info) Retrieve the contextual information of a device. Parameters dev_id The identifier of the device. dev_info A pointer to a structure of type <code>rte_cryptodev_info</code> to be filled with the contextual information of the device. Note: The capabilities field of <code>dev_info</code> is set to point to the first element of an array of struct <code>rte_cryptodev_capabilities</code> . The element after the last valid element has its <code>op</code> field set to <code>RTE_CRYPTOP_TYPE_UNDEFINED</code> .	void __rte_experimental rte_compressdev_info_get (uint8_t dev_id, struct rte_compressdev_info * dev_info) Retrieve the contextual information of a device. Parameters dev_id Compress device identifier dev_info A pointer to a structure of type <code>rte_compressdev_info</code> to be filled with the contextual information of the device. Note: The capabilities field of <code>dev_info</code> is set to point to the first element of an array of struct <code>rte_compressdev_capabilities</code> . The element after the last valid element has its <code>op</code> field set to <code>RTE_COMP_ALGO_LIST_END</code> . "line 147, 160"
	int __rte_experimental rte_compressdev_private_xform_create (uint8_t dev_id, const struct rte_comp_xform * xform, void ** private_xform) -- This should alloc a <code>private_xform</code> from the device's mempool and initialise it. The application should call this API when setting up for stateless processing on a device. If it returns non-shareable, then the appl cannot share this handle with multiple in-flight ops and should call this API again to get a separate handle for every in-flight op. The handle returned is only valid for use with ops of <code>op_type STATELESS</code> . Parameters dev_id Compress device identifier xform xform data private_xform Pointer to where PMD's <code>private_xform</code> handle should be stored Returns if successful returns 0 and valid <code>private_xform</code> handle <0 in error cases Returns <code>-EINVAL</code> if input parameters are invalid. Returns <code>-ENOTSUP</code> if comp device does not support the comp transform. Returns <code>-ENOMEM</code> if the <code>private_xform</code> could not be allocated. "line 150, 163, 220, 297"
static unsigned static unsigned rte_crypto_op_bulk_alloc (struct rte_mempool *mempool, enum rte_crypto_op_type type, struct rte_crypto_op **ops, uint16_t nb_ops) - Bulk	int __rte_experimental rte_comp_op_bulk_alloc (struct rte_mempool *mempool, struct rte_comp_op **ops, uint16_t nb_ops) - Bulk allocate operations from a mempool with default parameters set

<p>allocate crypto operations from a mempool with default parameters set</p> <p>Parameters Mempool crypto operation mempool type operation type to allocate ops Array to place allocated crypto operations nb_ops Number of crypto operations to allocate</p> <p>Returns nb_ops if the number of operations requested were allocated. 0 if the requested number of ops are not available. None are allocated in this case. (struct rte_mempool *mempool, enum rte_crypto_op_type type, struct rte_crypto_op **ops, uint16_t nb_ops) - Bulk allocate crypto operations from a mempool with default parameters set</p>	<p>Parameters mempool Compress operation mempool ops Array to place allocated operations nb_ops Number of operations to allocate</p> <p>Returns nb_ops: Success, the nb_ops requested was allocated 0: Not enough entries in the mempool; no ops are retrieved.</p> <p style="text-align: right;">”line 176, 179”</p>
<p>static int rte_pktmbuf_alloc_bulk (struct rte_mempool * pool, struct rte_mbuf ** mbufs, unsigned count) Allocate a bulk of mbufs, initialize refcnt and reset the fields to default values.</p> <p>Parameters pool The mempool from which mbufs are allocated. mbufs Array of pointers to mbufs count Array size</p> <p>Returns 0: Success -ENOENT: Not enough entries in the mempool; no mbufs are retrieved.</p> <p style="text-align: right;">”line 185, 188, 191, 194”</p>	
<p>static char * rte_pktmbuf_append (struct rte_mbuf *m, uint16_t len) - Append len bytes to an mbuf. Append len bytes to an mbuf and return a pointer to the start address of the added data. If there is not enough tailroom in the last segment, the function will return NULL, without modifying the mbuf.</p> <p>Parameters m The packet mbuf. len The amount of data to append (in bytes).</p> <p>Returns A pointer to the start of the newly appended data, or NULL if there is not enough tailroom space in the last segment</p> <p style="text-align: right;">”line 202, 209”</p>	
<p>rte_crypto_op_ctod_offset(c, t, o) ((t)((char *) (c) + (o))) A macro that points to an offset from the start of the crypto operation structure (rte_crypto_op)</p> <p>The returned pointer is cast to type t.</p> <p>Parameters c The crypto operation. o The offset from the start of the crypto operation. t The type to cast the result into.</p>	<p>rte_comp_op_ctod_offset(c, t, o) \ ((t)((char *) (c) + (o))) A macro that returns the physical address that points to an offset from the start of the comp operation (rte_comp_op)</p> <p>Parameters c The comp operation o The offset from the start of the comp operation to calculate address from (it is not used in the code)</p>

<p>The C library function void *memcpy(void *str1, const void *str2, size_t n) copies n characters from memory area str2 to memory area str1 void *memcpy(void *str1, const void *str2, size_t n)</p> <p>Parameters str1 – This is pointer to the destination array where the content is to be copied, type-casted to a pointer of type void*.</p> <p>str2 – This is pointer to the source of data to be copied, type-casted to a pointer of type void*.</p> <p>n – This is the number of bytes to be copied.</p> <p>Return Value This function returns a pointer to destination, which is str1.</p>	
<p>rte_pktmbuf_mtod(m, t) A macro that points to the start of the data in the mbuf. The returned pointer is cast to type t. Before using this function, the user must ensure that the first segment is large enough to accommodate its data.</p> <p>Parameters m The packet mbuf. t The type to cast the result into.</p> <p style="text-align: right;">”line 239, 282, 308, 340”</p>	
<p>rte_crypto_op Struct Reference Cryptographic Operation.</p> <p>This structure contains data relating to performing cryptographic operations. This operation structure is used to contain any operation which is supported by the cryptodev API, PMDs should check the type parameter to verify that the operation is a support function of the device. Crypto operations are enqueued and dequeued in crypto PMDs using the rte_cryptodev_enqueue_burst() / rte_cryptodev_dequeue_burst() .</p> <p>Data Fields struct rte_mempool * mempool rte_iova_t phys_addr union { struct rte_crypto_sym_op sym [0] struct rte_crypto_asym_op asym [0] }; uint8_t type uint8_t status uint8_t sess_type uint8_t reserved [3] uint16_t private_data_offset</p> <p>Explanation of Data Fields</p>	<p>rte_comp_op Struct Reference Compression Operation.</p> <p>This structure contains data relating to performing a compression operation on the referenced mbuf data buffers.</p> <p>Comp operations are enqueued and dequeued in comp PMDs using the rte_compressdev_enqueue_burst() / rte_compressdev_dequeue_burst() APIs.</p> <p>Data Fields struct rte_mempool * mempool enum type rte_iova_t iova_addr struct rte_mbuf * m_src struct rte_mbuf * m_dst enum rte_comp_flush_flag flush_flag uint64_t input_chksum uint64_t output_chksum uint32_t consumed uint32_t produced uint64_t debug_status uint8_t status void * private_xform void * stream uint32_t offset Note:</p>

<p>struct rte_mempool * mempool crypto operation mempool which operation is allocated from</p> <p>typedef uint64_t rte_iova_t physical address of crypto operation</p> <p>struct rte_crypto_sym_op sym [0] Symmetric operation parameters uint8_t type operation type</p> <p>operation uint8_t status -- status - this is reset to RTE_CRYPTOP_STATUS_NOT_PROCESSED on allocation from mempool and will be set to RTE_CRYPTOP_STATUS_SUCCESS after crypto operation is successfully processed by a crypto PMD</p> <p>uint16_t private_data_offset Offset to indicate start of private data (if any). The offset is counted from the start of the rte_crypto_op including IV. The private data may be used by the application to store information which should remain untouched in the library/driver</p> <p>rte_crypto_sym_op Struct Reference</p> <p>struct rte_mbuf * m_src Symmetric Cryptographic Operation.</p> <p>This structure contains data relating to performing symmetric cryptographic processing on a referenced mbuf data buffer.</p> <p>When a symmetric crypto operation is enqueued with the device for processing it must have a valid rte_mbuf structure attached, via m_src parameter, which contains the source data which the crypto operation is to be performed on. While the mbuf is in use by a crypto operation no part of the mbuf should be changed by the application as the device may read or write to any part of the mbuf. In the case of hardware crypto devices some or all of the mbuf may be DMAed in and out of the device, so writing over the original data, though only the part specified by the rte_crypto_sym_op for transformation will be changed. Out-of-place (OOP) operation, where the source mbuf is different to the destination mbuf, is a special case. Data will be copied from m_src to m_dst. The part copied includes all the parts of the source mbuf that will be operated on, based on the cipher.data.offset+cipher.data.length and auth.data.offset+auth.data.length values in the rte_crypto_sym_op. The part indicated by the cipher parameters will be transformed, any extra data around this indicated by the auth parameters will be</p>	<p>uint32_t length Note: uint8_t * digest</p> <p>NOTE:</p> <p>-op_type- Compression operation type: This field is not included in the official referral file. But it is used for real implementation.</p> <p>- The 'offset' field specified in the reference file is invalid, but src.offset / dst.offset is valid for actual implementation.</p> <p>-Also, 'length' specified in the reference doesn't work in actual implementation, but Also, src.length/ dst.length does work.</p> <p>Data Fields</p> <p>type- Compression operation type (Use op_type instead of this field.)</p> <p>Enumerator:</p> <p>-RTE_COMP_OP_STATELESS All data to be processed is submitted in the op, no state or history from previous ops is used and none will be stored for future ops. Flush flag must be set to either FLUSH_FULL or FLUSH_FINAL.</p> <p>-RTE_COMP_OP_STATEFUL There may be more data to be processed after this op, it's part of a stream of data. State and history from previous ops can be used and resulting state and history can be stored for future ops, depending on flush flag.</p> <p>struct rte_mempool * mempool --Pool from which operation is allocated</p> <p>rte_iova_t iova_add -- IOVA address of this operation, IO address of the buffer</p> <p>struct rte_mbuf * m_src -- source mbuf The total size of the input buffer(s) can be retrieved using rte_pktmbuf_pkt_len(m_src). The max data size which can fit in a single mbuf is limited by the uint16_t rte_mbuf.data_len to 64k-1. If the input data is bigger than this it can be passed to the PMD in a chain of mbufs if the PMD's capabilities indicate it supports this.</p> <p>struct rte_mbuf * m_dst -- destination mbuf The total size of the output buffer(s) can be retrieved using rte_pktmbuf_pkt_len(m_dst). The max data size which can fit in a single mbuf is limited by the uint16_t</p>
--	---

<p>copied unchanged from source to destination mbuf. Also in OOP operation the cipher.data.offset and auth.data.offset apply to both source and destination mbufs. As these offsets are relative to the data_off parameter in each mbuf this can result in the data written to the destination buffer being at a different alignment, relative to buffer start, to the data in the source buffer.</p> <pre> struct rte_mbuf * m_dst struct rte_cryptodev_sym_session * session struct rte_crypto_sym_xform * xform struct rte_security_session * sec_session uint32_t offset uint32_t length struct { uint32_t offset uint32_t length } data uint8_t * data rte_iova_t phys_addr struct { uint8_t * data rte_iova_t phys_addr } digest struct { uint8_t * data rte_iova_t phys_addr } aad struct { uint32_t offset uint32_t length } data struct { uint32_t offset uint32_t length } data struct { uint8_t * data rte_iova_t phys_addr } digest </pre> <p>uint32_t offset Starting point for AEAD processing, specified as number of bytes from start of packet in source buffer.</p> <p>Starting point for cipher processing, specified as number of bytes from start of data in the source buffer. The result of the cipher operation will be written back into the output buffer starting at this location.</p> <p>uint32_t length The message length, in bytes, of the source buffer on which the cryptographic operation will be computed. This must be a multiple of the block size</p>	<p>rte_mbuf.data_len to 64k-1. If the output data is expected to be bigger than this a chain of mbufs can be passed to the PMD if the PMD's capabilities indicate it supports this.</p> <p>enum rte_comp_flush_flag flush_flag -- Defines flush characteristics for the output data. Only applicable in compress direction</p> <p>uint64_t input_chksum -- An input checksum can be provided to generate a cumulative checksum across sequential blocks in a STATELESS stream. Checksum type is as specified in xform chksum_type.</p> <p>uint64_t output_chksum -- If a checksum is generated it will be written in here. Checksum type is as specified in xform chksum_type.</p> <p>uint32_t consumed -- The number of bytes from the source buffer which were compressed/decompressed.</p> <p>uint32_t produced -- The number of bytes written to the destination buffer which were compressed/decompressed.</p> <p>uint64_t debug_status -- Status of the operation is returned in the status param. This field allows the PMD to pass back extra pmd-specific debug information. Value is not defined on the API.</p> <p>uint8_t status -- Operation status - use values from enum rte_comp_status. This is reset to RTE_COMP_OP_STATUS_NOT_PROCESSED on allocation from mempool and will be set to RTE_COMP_OP_STATUS_SUCCESS after operation is successfully processed by a PMD.</p> <p>void * private_xform -- Stateless private PMD data derived from an rte_comp_xform. A handle returned by rte_compressdev_private_xform_create() must be attached to operations of op_type RTE_COMP_STATELESS.</p> <p>void * stream -- Private PMD data derived initially from an rte_comp_xform, which holds state and history data and evolves as operations are processed. rte_compressdev_stream_create() must be called on a device for all STATEFUL data streams and the resulting stream attached to the one or more operations associated with the data stream. All operations in a stream must be sent to the same device.</p> <p>uint32_t offset --src.offset in real code Starting point for compression or decompression, specified as number of bytes from</p>
---	--

<p>The message length, in bytes, of the source buffer on which the cryptographic operation will be computed. This must be a multiple of the block size if a block cipher is being used. This is also the same as the result length.</p>	<p>start of packet in source buffer. This offset starts from the first segment of the buffer, in case the m_src is a chain of mbufs. Starting point for checksum generation in compress direction.</p> <p>-dst.offset in real code-- Starting point for writing output data, specified as number of bytes from start of packet in dest buffer. This offset starts from the first segment of the buffer, in case the m_dst is a chain of mbufs. Starting point for checksum generation in decompress direction.</p> <p>uint32_t length -src.length in real code-- The length, in bytes, of the data in source buffer to be compressed or decompressed. Also the length of the data over which the checksum should be generated in compress direction</p> <p>uint8_t * digest -- Output buffer to store hash output, if enabled in xform. Buffer would contain valid value only after an op with flush flag = RTE_COMP_FLUSH_FULL/FLUSH_FINAL is processed successfully. Length of buffer should be contiguous and large enough to accommodate digest produced by specific hash algo.</p> <p style="text-align: right;">"line 173"</p>
Enqueue/Dequeue Burst APIs	
<p>static uint16_t rte_cryptodev_enqueue_burst (uint8_t dev_id, uint16_t qp_id, struct rte_crypto_op **ops, uint16_t nb_ops) Enqueue a burst of operations for processing on a crypto device.</p> <p>The rte_cryptodev_enqueue_burst() function is invoked to place crypto operations on the queue qp_id of the device designated by its dev_id.</p> <p>The nb_ops parameter is the number of operations to process which are supplied in the ops array of rte_crypto_op structures.</p> <p>The rte_cryptodev_enqueue_burst() function returns the number of operations it actually enqueued for processing. A return value equal to nb_ops means that all packets have been enqueued.</p> <p>Parameters dev_id The identifier of the device. qp_id The index of the queue pair which packets are to be enqueued for processing. The value must be in the range [0, nb_queue_pairs - 1] previously supplied to rte_cryptodev_configure.</p>	<p>uint16_t __rte_experimental rte_compressdev_enqueue_burst (uint8_t dev_id, uint16_t qp_id, struct rte_comp_op **ops, uint16_t nb_ops) Enqueue a burst of operations for processing on a compression device.</p> <p>The rte_compressdev_enqueue_burst() function is invoked to place comp operations on the queue qp_id of the device designated by its dev_id.</p> <p>The nb_ops parameter is the number of operations to process which are supplied in the ops array of rte_comp_op structures.</p> <p>The rte_compressdev_enqueue_burst() function returns the number of operations it actually enqueued for processing. A return value equal to nb_ops means that all packets have been enqueued.</p> <p>Note All compression operations are Out-of-place (OOP) operations, as the size of the output data is different to the size of the input data. The rte_comp_op contains both input and output parameters and is the vehicle for the application to pass data into and out of the PMD. While an op is inflight, i.e. once it has been enqueued, the private_xform or</p>

<p>ops The address of an array of nb_ops pointers to rte_crypto_op structures which contain the crypto operations to be processed.</p> <p>nb_ops The number of operations to process.</p> <p>Returns</p> <p>The number of operations actually enqueued on the crypto device. The return value can be less than the value of the nb_ops parameter when the crypto devices queue is full or if invalid parameters are specified in a rte_crypto_op.</p>	<p>stream attached to it and any mbufs or memory referenced by it should not be altered or freed by the application. The PMD may use or change some of this data at any time until it has been returned in a dequeue operation.</p> <p>The flush flag only applies to operations which return SUCCESS. In OUT_OF_SPACE cases whether STATEFUL or STATELESS, data in dest buffer is as if flush flag was FLUSH_NONE.</p> <p>flush flag only applies in compression direction. It has no meaning for decompression.</p> <p>: operation ordering is not maintained within the queue pair.</p> <p>Parameters</p> <p>dev_id Compress device identifier</p> <p>qp_id The index of the queue pair on which operations are to be enqueued for processing. The value must be in the range [0, nb_queue_pairs - 1] previously supplied to rte_compressdev_configure</p> <p>ops The address of an array of nb_ops pointers to rte_comp_op structures which contain the operations to be processed</p> <p>nb_ops The number of operations to process</p> <p>Returns</p> <p>The number of operations actually enqueued on the device. The return value can be less than the value of the nb_ops parameter when the comp devices queue is full or if invalid parameters are specified in a rte_comp_op.</p> <p style="text-align: right;">”line 254, 311”</p>
<p>static uint16_t rte_cryptodev_dequeue_burst (uint8_t dev_id, uint16_t qp_id, struct rte_crypto_op **ops, uint16_t nb_ops) Dequeue a burst of processed crypto operations from a queue on the crypto device. The dequeued operation are stored in rte_crypto_op structures whose pointers are supplied in the ops array.</p> <p>The rte_cryptodev_dequeue_burst() function returns the number of ops actually dequeued, which is the number of rte_crypto_op data structures effectively supplied into the ops array.</p> <p>A return value equal to nb_ops indicates that the queue contained at least nb_ops operations, and this is likely to signify that other processed operations remain in the devices output queue. Applications implementing a "retrieve as many processed operations as possible" policy can check this specific case and keep invoking the rte_cryptodev_dequeue_burst() function until a value less than nb_ops is returned.</p> <p>The rte_cryptodev_dequeue_burst() function does not provide any error notification to avoid the corresponding overhead.</p>	<p>uint16_t __rte_experimental</p> <p> rte_compressdev_dequeue_burst (uint8_t dev_id, uint16_t qp_id, struct rte_comp_op **ops, uint16_t nb_ops) Dequeue a burst of processed compression operations from a queue on the comp device. The dequeued operation are stored in rte_comp_op structures whose pointers are supplied in the ops array.</p> <p>The rte_compressdev_dequeue_burst() function returns the number of ops actually dequeued, which is the number of rte_comp_op data structures effectively supplied into the ops array.</p> <p>A return value equal to nb_ops indicates that the queue contained at least nb_ops operations, and this is likely to signify that other processed operations remain in the devices output queue. Applications implementing a "retrieve as many processed operations as possible" policy can check this specific case and keep invoking the rte_compressdev_dequeue_burst() function until a value less than nb_ops is returned.</p> <p>The rte_compressdev_dequeue_burst() function does not provide any error notification to avoid the corresponding overhead.</p>

<p>Parameters</p> <p>dev_id The symmetric crypto device identifier</p> <p>qp_id The index of the queue pair from which to retrieve processed packets. The value must be in the range [0, nb_queue_pair - 1] previously supplied to rte_cryptodev_configure().</p> <p>ops The address of an array of pointers to rte_crypto_op structures that must be large enough to store nb_ops pointers in it.</p> <p>nb_ops The maximum number of operations to dequeue.</p> <p>Returns</p> <p>The number of operations actually dequeued, which is the number of pointers to rte_crypto_op structures effectively supplied to the ops array.</p>	<p>Note</p> <p>: operation ordering is not maintained within the queue pair.</p> <p>: In case op status = OUT_OF_SPACE_TERMINATED, op.consumed=0 and the op must be resubmitted with the same input data and a larger output buffer. op.produced is usually 0, but in decompression cases a PMD may return > 0 and the application may find it useful to inspect that data. This status is only returned on STATELESS ops.</p> <p>: In case op status = OUT_OF_SPACE_RECOVERABLE, op.produced can be used and next op in stream should continue on from op.consumed+1 with a fresh output buffer. Consumed=0, produced=0 is an unusual but allowed case. There may be useful state/history stored in the PMD, even though no output was produced yet.</p> <p>Parameters</p> <p>dev_id Compress device identifier</p> <p>qp_id The index of the queue pair from which to retrieve processed operations. The value must be in the range [0, nb_queue_pair - 1] previously supplied to rte_compressdev_configure()</p> <p>ops The address of an array of pointers to rte_comp_op structures that must be large enough to store nb_ops pointers in it</p> <p>nb_ops The maximum number of operations to dequeue</p> <p>Returns</p> <p>The number of operations actually dequeued, which is the number of pointers to rte_comp_op structures effectively supplied to the ops array.</p> <p style="text-align: right;">”line 273, 329”</p>
--	---

Table1: Functions and structs of OpenSSL and Zlib PMDs

5. Compression application

There are two compression PMDs: ZLIB and ISA-L. We developed a compression program using ZLIB PMD and described the general order of the application as follows:

1. **Preparing environment for deployment:** Installing and initialazation of DPDK and ZLIB compress PMD.
2. **Creating compress devices and memories**
 - A. Creating memory pools for packet and compress operations
 - B. Creating compress and decompress devices
 - C. Configuring devices
 - D. Configuring queue pair on the devices

- E. Starting compress devices.
- F. Defining compress and decompress xforms.
- G. Allocating memory to compress operation and processing
- H. Allocating memory to packets.

3. Setup and process compression operations

- A. Defining compress/decompress parameters
- B. Attaching compress xform to operation
- C. Enqueueing compress operation on the compress device for processing
- D. Dequeueing processed operation (compressed data) from the compress device.
- E. Attaching decompress xform to operation
- F. Enqueueing compressed data on the decompress device.
- G. Dequeueing uncompressed data/processed data from the decompress device.

6. Conclusion

In this technical report, we first described the general overview of DPDK, DPDK libraries and PMDs. Then we compared structs and functions of the applications using OpenSSL crypto and ZLIB compression PMD in Table 1. When programming this compress application, we found some differences in the `rte_comp_op` struct fields between DPDK's official reference file and actual implementation. For example, an `op_type` field is not specified in DPDK's official reference file, but this field is valid for actual implementation. In addition, the reference file contains only "offset" and "long" fields in `rte_comp_op` struct, but 'src.offset / dst.offset' and 'src.length / dst.length' were valid in the real code. We noted this kind of differences in Table 1. Finally, we described the general order of our compressed application. Appendix 1 and 2 include the DPDK, PMDs installation instructions, and compression source code for the PMD program.

Acknowledgment

I would like to thank the European Union, and the European Social Fund (EFOP-3.6.2-16-2017-00013, Thematic Fundamental Research Collaborations Grounding Innovation in Informatics and Infocommunications) for funding this work. I would also like to thank my supervisor for a great deal of valuable advice.

Reference

1. Dominik Scholz(2014), A Look at Intel's Dataplane Development Kit
2. Intel DPDK: Programmers guide, DPDK 19.02 and 19.05
3. Intel DPDK: Packet Processing on Intel Architecture, Presentation slides, 2012
4. Test application of compress PMD:
http://git.dpdk.org/dpdk/plain/app/test/test_compressdev.c
5. Andrej Yemelianov(2016), Introduction to DPDK: Architecture and Principles
(<https://blog.selectel.com/introduction-dpdk-architecture-principles/>)

Appendix 1: DPDK, crypto OpenSSL PMD, and compression ZLIB PMD installation guide

1. Download DPDK

```
wget https://fast.dpdk.org/rel/dpdk-19.02.tar.xz
tar xJf dpdk-<version>.tar.xz
```

2. Go to DPDK directory

```
cd dpdk-<version>
```

3. Install the following prerequisite packages

```
apt-get install libpcap-dev gcc make hugepages nim
apt-get install linux-headers-generic
apt-get install libnuma-dev
apt-get install libelf-dev
```

3. Install ZLIB and OpenSSL library

```
apt-get install zlib1g-dev
apt-get install libssl-dev
```

4. In order to enable this virtual compression PMD and crypto OpenSSL, user must:

Set CONFIG_RTE_LIBRTE_PMD_ZLIB=y in config/common_base.
Set CONFIG_RTE_LIBRTE_PMD_OPENSSL=y in config/common_base.

5. Build and install dpdk

```
make config T=x86_64-native-linuxapp-gcc
make
make install
```

Appendix 2: Source code of application using ZLIB compression PMD.

```
1 #include <rte_compressdev.h>
2 #include <unistd.h>
3 #include <rte_bus_vdev.h>
4
5 #define DEFAULT_WINDOW_SIZE 15
6 #define DEQUEUE_WAIT_TIME 10000
7 #define NUM_MAX_XFORMS 16
8 #define NUM_MAX_INFLIGHT_OPS 512
9 #define NUM_MBUFS 8192
10 #define POOL_CACHE_SIZE 256
11 #define NUM_OPS 1
12 #define BUFFER_SIZE 32
13 #define PRIV_SIZE 16
14
15 void print_bytes(unsigned char *, size_t);
16 void print_bytes(unsigned char *c, size_t n)
```

```

17  {
18  size_t i;
19  for(i = 0; i < n; i++)
20      printf("%02X ", c[i]);
21  printf("\n");
22  }
23 int main(int argc, char **argv)
24  {
25  srand ((unsigned int) time (NULL));
26  int ret;
27
28  /* Initialize EAL. */
29  ret = rte_eal_init(argc, argv);
30  if (ret < 0)
31      rte_exit(EXIT_FAILURE, "Invalid EAL arguments\n");
32  uint8_t socket_id = rte_socket_id();
33  struct rte_mempool *mbuf_pool, *op_pool;
34
35  /* Create the mbuf pool. */
36  mbuf_pool = rte_pktmbuf_pool_create("mbuf_pool",
37                                     NUM_MBUEFS,
38                                     POOL_CACHE_SIZE,
39                                     0,
40                                     RTE_MBUF_DEFAULT_BUF_SIZE,
41                                     socket_id);
42  if (mbuf_pool == NULL)
43      rte_exit(EXIT_FAILURE, "Cannot create mbuf pool\n");
44  unsigned int comp_op_private_data = PRIV_SIZE;
45
46  /* Create compression operation pool. */
47  op_pool = rte_comp_op_pool_create( "op_pool",
48                                     NUM_MBUEFS,
49                                     POOL_CACHE_SIZE,
50                                     comp_op_private_data,
51                                     socket_id);
52  if (op_pool == NULL)
53      rte_exit(EXIT_FAILURE, "Cannot create crypto op pool\n");
54
55  /** Compress device *****/
56  char args[128];
57  const char *compress_name = "compress_zlib";
58  snprintf(args, sizeof(args), "socket_id=%d", socket_id);
59  ret = rte_vdev_init(compress_name, args);
60  if (ret != 0)
61      rte_exit(EXIT_FAILURE, "Cannot create virtual device");
62  uint8_t cdev_id_compress = rte_compressdev_get_dev_id(compress_name);
63
64  /** Decompress device *****/
65  const char *compress_name2 = "compress_zlib1";
66  snprintf(args, sizeof(args), "socket_id=%d", socket_id);
67  ret = rte_vdev_init(compress_name2, args);
68  if (ret != 0)
69      rte_exit(EXIT_FAILURE, "Cannot create virtual device");
70  uint8_t cdev_id_decompress = rte_compressdev_get_dev_id(compress_name2);
71
72  /* Configure the compress device. Allocation of resources */

```

```

73 struct rte_compressdev_config conf = { //structure is used to pass configuration parameters.
74     .socket_id = rte_socket_id(),
75     .nb_queue_pairs = 1,
76     .max_nb_priv_xforms = NUM_MAX_XFORMS,
77     .max_nb_streams = 1
78 };
79
80 if (rte_compressdev_configure(cdev_id_compress, &conf) < 0) // rte_compressdev_configure API is used to
configure a compression device
81     rte_exit(EXIT_FAILURE, "Failed to configure compressdev %u", cdev_id_compress);
82
83 if (rte_compressdev_configure(cdev_id_compress, &conf) == 0)
84     printf( "Compress device is succesfully configured. \n");
85
86 if (rte_compressdev_queue_pair_setup(cdev_id_compress, 0, NUM_MAX_INFLIGHT_OPS,
87     rte_socket_id()) < 0)
88     rte_exit(EXIT_FAILURE, "Failed to setup queue pair\n");
89
90 if (rte_compressdev_queue_pair_setup(cdev_id_compress, 0, NUM_MAX_INFLIGHT_OPS,
91     rte_socket_id()) == 0)
92     printf("The queue pair is successfully configured in the compress device. \n");
93
94 if (rte_compressdev_start(cdev_id_compress) < 0)
95     rte_exit(EXIT_FAILURE, "Failed to start device\n");
96 else
97     printf("Compress device successfully created \n");
98
99 /* Configure the decompress device. Allocation of resources */
100 if (rte_compressdev_configure(cdev_id_decompress, &conf) < 0)
101     rte_exit(EXIT_FAILURE, "Failed to configure compressdev %u", cdev_id_decompress);
102
103 if (rte_compressdev_configure(cdev_id_decompress, &conf) == 0)
104     printf( "Decompress device is succesfully configured. \n");
105
106 if (rte_compressdev_queue_pair_setup(cdev_id_decompress, 0, NUM_MAX_INFLIGHT_OPS,
107     rte_socket_id()) < 0)
108     rte_exit(EXIT_FAILURE, "Failed to setup queue pair\n");
109
110 if (rte_compressdev_queue_pair_setup(cdev_id_decompress, 0, NUM_MAX_INFLIGHT_OPS,
111     rte_socket_id()) == 0)
112     printf("The queue pair is successfully configured in the decompress device. \n");
113
114 if (rte_compressdev_start(cdev_id_decompress) < 0)
115     rte_exit(EXIT_FAILURE, "Failed to start device\n");
116 else
117     printf("Decompress device successfully created \n");
118
119 /* Setup compress transform */
120 struct rte_comp_xform compress_xform = {
121     .type = RTE_COMP_COMPRESS,
122     .compress = {
123         .algo = RTE_COMP_ALGO_DEFLATE,
124         .deflate = {
125             .huffman = RTE_COMP_HUFFMAN_DEFAULT
126         },
127         .level = RTE_COMP_LEVEL_PMD_DEFAULT,

```



```

128     .chksum = RTE_COMP_CHECKSUM_NONE,
129     .window_size = DEFAULT_WINDOW_SIZE,
130     .hash_algo = RTE_COMP_HASH_ALGO_NONE
131 }
132 };
133
134 /* Setup decompress transform */
135 struct rte_comp_xform decompress_xform = {
136     .type = RTE_COMP_DECOMPRESS,
137     .decompress = {
138         .algo = RTE_COMP_ALGO_DEFLATE,
139         .chksum = RTE_COMP_CHECKSUM_NONE,
140         .window_size = DEFAULT_WINDOW_SIZE,
141     }
142 };
143
144 struct rte_compressdev_info dev_info;
145 void *priv_xform = NULL;
146 int shareable;
147 rte_compressdev_info_get(cdev_id_compress, &dev_info);
148 if(dev_info.capabilities->comp_feature_flags & RTE_COMP_FF_SHAREABLE_PRIV_XFORM)
149     {
150     rte_compressdev_private_xform_create(cdev_id_compress, &compress_xform, &priv_xform);
151     printf("success");
152     }
153 else
154     {
155     shareable = 0;
156     printf("Non-shareable is configured in the compress device. \n");
157     }
158 struct rte_compressdev_info dev_info2;
159 void *priv_xform2 = NULL;
160 rte_compressdev_info_get(cdev_id_decompress, &dev_info2);
161 if(dev_info2.capabilities->comp_feature_flags & RTE_COMP_FF_SHAREABLE_PRIV_XFORM)
162     {
163     rte_compressdev_private_xform_create(cdev_id_decompress, &decompress_xform, &priv_xform2);
164     printf("success");
165     }
166 else
167     {
168     shareable = 0;
169     printf("Non-shareable is configured in the decompress device. \n");
170     }
171
172 /* Create an op pool and allocate ops */
173 struct rte_comp_op *comp_ops[NUM_OPS], *processed_ops[NUM_OPS] ; /*rte_comp_op structure contains
data relating to performing
174
a compression operation on
the referenced mbuf data buffers.*/
175
176 if (rte_comp_op_bulk_alloc(op_pool, comp_ops, NUM_OPS) == 0)
177     rte_exit(EXIT_FAILURE, "Not enough crypto operations available\n");
178
179 if (rte_comp_op_bulk_alloc(op_pool, processed_ops, NUM_OPS) == 0)
180     rte_exit(EXIT_FAILURE, "Not enough crypto operations available\n");
181

```

```

182 /* Prepare source and destination mbufs for compression operations*/
183 struct rte_mbuf *comp_mbufs[BUFFER_SIZE], *uncomp_mbufs[BUFFER_SIZE];
184
185 if (rte_pktmbuf_alloc_bulk(mbuf_pool, comp_mbufs, NUM_OPS) < 0)
186     rte_exit(EXIT_FAILURE, "Not enough room in the source mbuf");
187
188 if (rte_pktmbuf_alloc_bulk(mbuf_pool, comp_mbufs, NUM_OPS) == 0)
189     printf("Allocate a source bulk of comp_mbufs, initialize refcnt and reset the fields to default values\n");
190
191 if (rte_pktmbuf_alloc_bulk(mbuf_pool, uncomp_mbufs, NUM_OPS) < 0)
192     rte_exit(EXIT_FAILURE, "Not enough room in the dst mbuf");
193
194 if (rte_pktmbuf_alloc_bulk(mbuf_pool, uncomp_mbufs, NUM_OPS) == 0)
195     printf("Allocate a destination bulk of mbufs, initialize refcnt and reset the fields to default values\n");
196
197 /* Prepare source and destination mbufs for compression operations*/
198 unsigned int i;
199
200 for (i = 0; i < NUM_OPS; i++)
201     {
202         if (rte_pktmbuf_append(uncomp_mbufs[i], BUFFER_SIZE) == NULL)
203             rte_exit(EXIT_FAILURE, "Not enough room in the mbuf\n");
204         comp_ops[i]->m_src=uncomp_mbufs[i];
205     }
206
207 for (i = 0; i < NUM_OPS; i++)
208     {
209         if (rte_pktmbuf_append(comp_mbufs[i], BUFFER_SIZE) == NULL)
210             rte_exit(EXIT_FAILURE, "Not enough room in the mbuf\n");
211         comp_ops[i]->m_dst=comp_mbufs[i];
212     }
213 uint16_t num_enqd, num_deqd;
214
215 // Set up the compress operations.
216 for (i = 0; i < NUM_OPS; i++)
217     {
218         if (!shareable)
219             {
220                 if(rte_compressdev_private_xform_create(cdev_id_compress, &compress_xform, &comp_ops[i]-
221 >private_xform)==0)
222                 {
223                     printf("Compresss xform created in compress device \n");
224                 }
225             }
226         else
227             {
228                 comp_ops[i]->private_xform = priv_xform;
229             }
230
231         comp_ops[i]->op_type = RTE_COMP_OP_STATELESS;
232         comp_ops[i]->flush_flag = RTE_COMP_FLUSH_FINAL;
233         comp_ops[i]->src.offset = 0;
234         comp_ops[i]->dst.offset = 0;
235         comp_ops[i]->src.length = BUFFER_SIZE;
236         printf("Original data: ");

```

```

237     int j;
238     struct rte_mbuf *m= comp_ops[i]->m_src;
239     unsigned char *data = rte_pktmbuf_mtod(m, unsigned char*); //A macro that points to the start of the data
in the mbuf.
240     for (j = 1; j < 16; j++)
241         data[j] = 9;
242     comp_ops[i]->input_chksm = 0;
243     print_bytes(data, 16);
244
245     /*rte_compressdev_enqueue_burst---Enqueue a burst of operations for processing on a compression device.
246     *The rte_comp_op contains both input and output parameters and is the
247     * vehicle for the application to pass data into and out of the PMD. While an
248     * op is inflight, i.e. once it has been enqueued, the private_xform
249     * attached to it and any mbufs or memory referenced by it should not be altered
250     * or freed by the application. The PMD may use or change some of this data at
251     * any time until it has been returned in a dequeue operation.
252     *
253     */
254     num_enqd = rte_compressdev_enqueue_burst(cdev_id_compress, 0, comp_ops , NUM_OPS);
255     if(comp_ops[i]->status == RTE_COMP_OP_STATUS_SUCCESS)
256     {
257         printf("Enqueue operation completed succesfully in compress device  \n");
258     }
259     printf("comp_ops[i]->consumed \n");
260     printf("%d\n", comp_ops[i]->consumed);
261     printf("comp_ops[i]->produced \n");
262     printf("%d\n", comp_ops[i]->produced);
263     printf("processed_ops[i]->consumed \n");
264     printf("%d\n", processed_ops[i]->consumed);
265     printf(" processed_ops[i]->produced \n");
266     printf("%d\n", processed_ops[i]->produced);
267     }
268
269     uint16_t total_processed_ops=0;
270     usleep(DEQUEUE_WAIT_TIME);
271     do {
272         /*Dequeue a burst of processed compression operations from a queue on the compress device.*/
273         num_deqd = rte_compressdev_dequeue_burst(cdev_id_compress, 0 , processed_ops, NUM_OPS);
274         total_processed_ops += num_deqd;
275         /* Check if operation was processed successfully */
276         for (i = 0; i < num_deqd; i++) {
277             if (processed_ops[i]->status != RTE_COMP_OP_STATUS_SUCCESS)
278                 rte_exit(EXIT_FAILURE,"Some operations were not processed correctly");
279             else {
280                 printf("Dequeue operation completed successfully in compress device \n");
281                 struct rte_mbuf *m = processed_ops[i]->m_dst;
282                 unsigned char *data = rte_pktmbuf_mtod(m, unsigned char*);
283                 printf("Compressed data: ");
284                 print_bytes(data, 16);
285             }
286         }
287
288         rte_mempool_put_bulk(op_pool, (void **)processed_ops,num_deqd);
289     }
290     while (total_processed_ops < num_enqd); //push next op
291

```

```

292 //-----
293 for (i = 0; i < NUM_OPS; i++) {
294     if (!shareable)
295     {
296         /*rte_compressdev_private_xform_create should alloc a private_xform from the device's mempool and initialise
it.*/
297         rte_compressdev_private_xform_create(cdev_id_decompress, &decompress_xform, &comp_ops[i]-
>private_xform);
298         printf("Decompress xform created in decompress device \n");
299     }
300 else
301     {
302         comp_ops[i]->private_xform = priv_xform2;
303     }
304     printf("Compressed data in decompress device: ");
305     comp_ops[i]->m_src=processed_ops[i]->m_dst;
306     comp_ops[i]->m_dst=uncomp_mbufs[i];
307     struct rte_mbuf *m= comp_ops[i]->m_src;
308     unsigned char *data = rte_pktmbuf_mtod(m, unsigned char*);
309     print_bytes(data, 16);
310
311     num_enqd = rte_compressdev_enqueue_burst(cdev_id_decompress, 0, comp_ops , NUM_OPS);
312     if(comp_ops[i]->status == RTE_COMP_OP_STATUS_SUCCESS)
313     {
314         printf("Enqueue operation completed succesfully in decompress device \n");
315     }
316     printf("comp_ops[i]->consumed \n");
317     printf("%d\n", comp_ops[i]->consumed);
318     printf("comp_ops[i]->produced \n");
319     printf("%d\n", comp_ops[i]->produced);
320     printf("processed_ops[i]->consumed \n");
321     printf("%d\n", processed_ops[i]->consumed);
322     printf(" processed_ops[i]->produced \n");
323     printf("%d\n", processed_ops[i]->produced);
324     }
325     total_processed_ops=0;
326     usleep(DEQUEUE_WAIT_TIME);
327     do {
328         /*Dequeue a burst of processed compression operations from a queue on the compress device.*/
329         num_deqd = rte_compressdev_dequeue_burst(cdev_id_decompress, 0 , processed_ops, NUM_OPS);
330         total_processed_ops += num_deqd;
331         /* Check if operation was processed successfully */
332         for (i = 0; i < num_deqd; i++) {
333             if (processed_ops[i]->status != RTE_COMP_OP_STATUS_SUCCESS)
334                 rte_exit(EXIT_FAILURE,
335                     "Some operations were not processed correctly");
336         else
337             { printf("Dequeue operation completed succesfully in decompress device \n");
338
339                 struct rte_mbuf *m = processed_ops[i]->m_dst;
340                 unsigned char *data = rte_pktmbuf_mtod(m, unsigned char*);
341                 printf("Original data in decompress device: ");
342                 print_bytes(data, 16);
343             }
344         }
345     }

```

```
346 }  
347 while (total_processed_ops < num_enqd);  
348 }
```